



Introducing MongoDB

Assessable Learning Outcomes

- ▶ What is MongoDB?
- ▶ Why is it useful?
- ▶ How can I
 - ▶ Use a database, containing collections
 - ▶ Query existing collections
- ▶ What are the design principles?
 - ▶ Design a 'schema' for use in an application
- ▶ Choose a design and a type of database.

Document Database

- ▶ A document is a data structure composed of field and value pairs.
- ▶ MongoDB documents are similar to JSON objects.
- ▶ The values of fields may include other documents, arrays, and arrays of documents.

MongoDB advantages

- ▶ MongoDB says the advantages of using documents are:
 - ▶ Documents (i.e. objects) correspond to native data types in many programming languages.
 - ▶ Embedded documents and arrays reduce need for expensive joins.
 - ▶ Dynamic schema supports fluent polymorphism.

Key Features

- ▶ High Performance
 - ▶ Supports embedded data models (reduces I/O)
 - ▶ Allows Indexes
- ▶ Rich Query Language
 - ▶ Read, write, aggregation, text search, geospatial.
- ▶ High Availability
 - ▶ Uses replication
- ▶ Horizontally scalable through sharding.
- ▶ Supports heterogeneous storage engines.

CAP Theorem

- ▶ C – Consistency

- ▶ This is a given in a relational database
- ▶ If a RDBMS is not in a consistent state, it is not available.
 - ▶ Remember what happens when two sessions are accessing the same data – only the consistent data is available to both.

- ▶ A – Availability

- ▶ A MongoDB database will always have data available because it has replicas of every piece of data and it only promises 'eventual consistency'

- ▶ P – Partition tolerance (*tolerating network breaks*)

- ▶ In a distributed situation, RDBMS will not allow a transaction to go through unless all nodes involved are in a consistent state.
- ▶ In a distributed situation, MongoDB will go ahead with transactions for available nodes and let the others catch up later.

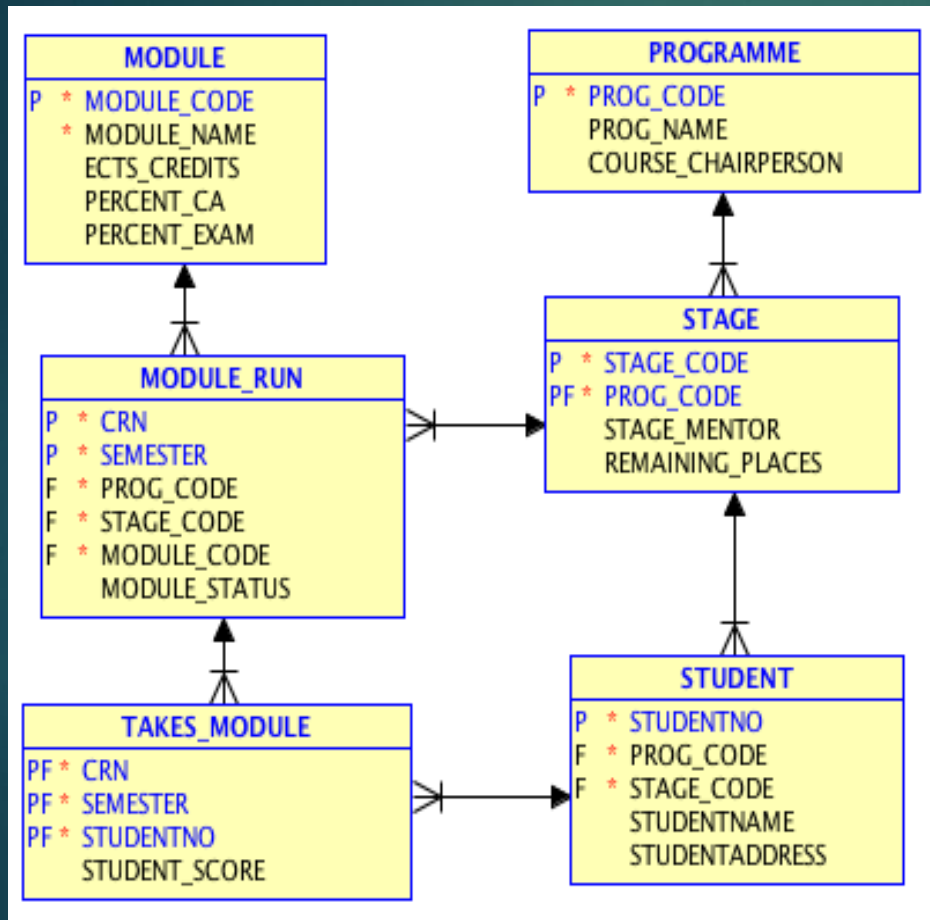
Using MongoDB

- ▶ We will experiment on a service that is hosted by DIT
- ▶ You will install your own version, on your laptop.

Remember

- ▶ There are no joins in a MongoDB query.
 - ▶ Queries may be piped, to result in the equivalent of a join.
- ▶ Transactions in MongoDB involving more than one instruction are not atomic.

Our student-modules ERD



► Note the ERD

- There are a lot of weak entities
- To get a transcript, a lot of joins are required.
- The structure is not intuitive.

Scenarios in MongoDB

- ▶ Add a student collection, embedding the module codes in the student.
- ▶ Student
 - ▶ Studentno (can be used as `_id`)
 - ▶ StudentName
 - ▶ Prog_code
 - ▶ Stage_code
 - ▶ Modules:
 - ▶ Module_code
 - ▶ Student_score

Scenarios in MongoDB

▶ Student

- ▶ Studentno (as _id)
- ▶ StudentName
- ▶ Prog_code
- ▶ Stage_code
- ▶ Modules:
 - ▶ Module_code
 - ▶ Student_score

```
▶ db.student.insert({ _id :  
  'C22345678',  
  StudentName: 'Joe Bloggs',  
  prog_code: 'DT222',  
  stage_code: 3,  
  modules: [  
    {Module_code: 'CMPU3010',  
     Student_score: 87},  
    {Module_code: 'CMPU3048',  
     Student_score: 62}])
```

```
db.student.insert([{\n  _id: 'C12345678',\n  studentname: 'Joe Bloggs',\n  prog_code: 'DT228',\n  stage_code: 3,\n  modules: {\n    {module_code: 'CMPU3010', student_score: 66},\n    {module_code: 'CMPU3047', student_score: 45}\n  },\n  {\n    _id: 'C12345679',\n    studentname: 'Jane Bloggs',\n    prog_code: 'DT228',\n    stage_code: 4},\n  {\n    _id: 'C12345670',\n    studentname: 'Enda Kenny',\n    prog_code: 'DT222',\n    prog_year: 1,\n    societies_joined: ['Young Fine Gael', 'Rowing']\n  }\n])
```

Sample entry

► Note:

- We don't need to leave space for null attributes
- We can add extra fields whenever we like.

```
db.student.insert([{
  _id: 'C12345678',
  studentname: 'Joe Bloggs',
  prog_code: 'DT228',
  stage_code: 3,
  modules: [
    {module_code: 'CMPU3010', student_score: 66},
    {module_code: 'CMPU3047', student_score: 45}
  ],
  {
    _id: 'C12345679',
    studentname: 'Jane Bloggs',
    prog_code: 'DT228',
    stage_code: 4},
    {
      _id: 'C12345670',
      studentname: 'Enda Kenny',
      prog_code: 'DT222',
      prog_year: 1,
      societies_joined: ['Young Fine Gael', 'Rowing']}
  ]
})
```

Sample entry

► Note:

- We don't need to leave space for null attributes
 - We can add extra fields whenever we like.
- MongoDB will not stop you from calling an attribute different names in different documents, but this hampers searching.

Or this way...

- ▶ Add a module collection, embedding the students who sit it in the module:
- ▶ Module
 - ▶ Modulecode
 - ▶ Title
 - ▶ ECTS credits
 - ▶ CA percent
 - ▶ Exam percent
 - ▶ Students:
 - ▶ StudentNo

Adding modules

- ▶ Module
 - ▶ Modulecode
 - ▶ Title
 - ▶ ECTS credits
 - ▶ CA percent
 - ▶ Exam percent
 - ▶ Students:
 - ▶ StudentNo

```
db.module.insert([{
    _id: 'CMPU3010',
    title: 'Databases 2',
    ECTS: 5,
    CAPercent: 40,
    ExamPercent: 60,
    Students: ['C12345678', 'C12345671']
},
{
    _id: 'CMPU1011',
    title: 'Build a PC',
    ECTS: 5,
    CAPercent: 100,
    Students:
    ['C16123456', 'C16123457', 'C16123458']]})
```

Or this way

▶ Module

- ▶ Modulecode
- ▶ Title
- ▶ ECTS credits
- ▶ CA percent
- ▶ Exam percent
- ▶ Students:
 - ▶ StudentNo (reference)

▶ Student

- ▶ Studentno (can be used as _id)
- ▶ StudentName
- ▶ Prog_code
- ▶ Stage_code
- ▶ Modules:
 - ▶ Module_code (reference)
 - ▶ Student_score

Indexes in MongoDB

- ▶ As for relational databases, MongoDB indexes support the efficient execution of queries. Without indexes, MongoDB must perform a *collection scan* (the equivalent of a full-table scan in RDBMS).
- ▶ Indexes store a small portion of the collection's data set in an easy to traverse form. The index stores the value of a specific field or set of fields, ordered by the value of the field.

Indexes

- ▶ Every document is automatically indexed on the `_id` attribute at insertion.
- ▶ Indexes can be created for other attributes using the `db.collection.createIndex ()` method.



Getting started

CAN WE REUSE THE DATA WE ALREADY HAVE?

There are ways to reuse data

- ▶ By hand
 - ▶ Type the code into notepad++
- ▶ Very simple (no embedding)
 - ▶ Export your SQL output to an excel file. Find a converter online and run it.
 - ▶ Write a SQL query to generate it
- ▶ A little harder (with embedding)
 - ▶ Write a PL/SQL program to do it or
 - ▶ Use XML SQL commands to generate it.

Write a SQL query:

- ▶ Spool redirects output to wherever it is sent, but the output needs a bit of work when it's done.
- ▶ `spool "C:\wherever\students.js"`
- ▶ `SELECT 'db.student.insert({_id: '||studentno||', studentname: '||studentname||', prog_code: '||prog_code||', stage_code: '||stage_code||'})' FROM STUDENT;`
- ▶ `spool off;`

PL/SQL generation

- ▶ Use the spool feature again but run a program.
- ▶ The program needs 2 cursors
- ▶ For each student (use a student cursor):
 - ▶ Start the db..insert command
 - ▶ Use dbms_output.put (not put_line) to add attributes to it.
 - ▶ If the student takes modules , add a modules array:
 - ▶ For each module the student takes (new cursor)
 - ▶ Add the module code
 - ▶ Finish off the array
 - ▶ Close the cursor
 - ▶ Finish the insert command.

Additional material

- ▶ The MongoDB tutorial will get you up and running.
- ▶ There are some scripts that I have written / generated:
 - ▶ Studentxml.sql uses SQLXML to generate an XML file for conversion to JSON.
 - ▶ Students.xml is generated using SQL for XML. It can be easily converted to JSON.
 - ▶ Createflatstudentsql.sql uses SQL to generate students with no embedding
 - ▶ Createflatstudent.sql uses PL/SQL to generate students with no embedding
 - ▶ Createmongostudent.sql uses PL/SQL to generate students with embedded array of modules taken.